

Breaking SHA-256 with Physics

Patch-Overlap Constraint Search on the Observer Screen

Bernhard Mueller and the Pragma Research Team

May 10, 2026

Paper release: r1453 **Released:** May 18, 2026

Abstract

SHA-256 was designed to make search look stupid. In the black-box model, a silicon miner that makes Q classical hash queries against a target fraction p succeeds with probability at most Qp . That bound is the proof behind the brute-force security of ordinary silicon search. There is no slope to climb, no half-open door, no helpful smell of the right nonce. Full SHA-256d proof of work is brute force for digital hardware, and modern ASICs win by paying the same bill faster.

Our attack changes the physical layer on which the search is performed. Observer Patch Holography (OPH) starts with the physics of records: overlap, syndrome, repair, and fixed-point agreement [31]. Ordinary quantum mechanics appears as the record calculus seen after patches agree, while the search machine here is built on the layer underneath that calculus. SHA-256d mining then becomes a finite physical constraint body. The nonce, schedule, compression trajectory, and target face become four observer patches, and their shared collars become physical interfaces. A wrong local assignment emits a syndrome and dies where its consequences stop matching. A surviving mode is a record of global agreement.

In hardware, the map is physical. Bits become wave signs, Boolean relations become local energies, and nonlinear ports supply the carry and target gates that pure propagation cannot provide by itself. The exact digital hash stays outside the chamber as the referee. The mathematical layer is complete: exact Boolean-to-energy equivalence, four-patch collar consistency, and the physical-oracle theorem define the search speedup. The empirical layer is complete at prototype scale: the OPH/Karma optical path has produced pool-accepted kHeavyHash shares under exact digital verification. The open surface is engineering: cavity geometry, port nonlinearity, calibration stability, and SHA-256d scale-up. The SHA-256d construction replaces the Keccak/kHeavyHash core with the four-patch SHA-256d body, keeps the exact verifier, and measures speed by physical settling time and candidate enrichment. The measured and projected wave-overlap regimes give speedups of several orders of magnitude over conventional verifier rates, with the ideal endpoint being one zero-energy settling event per valid nonce.

Keywords: SHA-256; proof of work; Kaspas; kHeavyHash; physical cryptanalysis; optical computing; analog computing; Grover search; Observer Patch Holography; constraint satisfaction

1 Introduction

SHA-256 is a standard cryptographic hash function specified by NIST in FIPS 180-4 [1]. Bitcoin proof of work applies SHA-256 twice to an 80-byte block header and asks for an output below a target [3]. In the ideal-hash abstraction, each tested header is an independent draw from $\{0, 1\}^{256}$. If the target requires k leading zero bits, an unbiased miner succeeds with probability 2^{-k} per candidate and needs 2^k trials in expectation.

The ordinary attack surface is simple. A digital miner evaluates the whole compression pipeline for one candidate, then another, then another. Modern ASICs reduce the cost per candidate while preserving the mathematical shape of the search. Published cryptanalytic attacks reach reduced-step variants or related settings. Full SHA-256 preimage search remains a practical enumeration problem [11, 12, 13]. Quantum computers change the query count by Grover amplification: quadratically in the number of candidates. Published resource estimates for cryptographic SHA-256 preimage search remain far beyond fault-tolerant hardware [6, 8, 9, 10].

Observer Patch Holography suggests a different representation of the problem. The blog essay “P = NP on the Observer Screen” phrases the computational primitive in five words: “Verification is the force” [22]. The technical OPH papers develop the same idea through patch algebras, overlap consistency, syndromes, repair moves, and fixed-point consensus [33, 34, 32]. For SHA-256d mining, the primitive becomes an auditable patch network.

SHA-256d translates into a finite constraint network. The network has four patches:

- Patch A owns the nonce-bearing message tail, padding, length field, and the initial schedule words $W_0, \dots, W_{15}$.
- Patch B owns the message-schedule recurrence

$$W_t = \sigma_1(W_{t-2}) + W_{t-7} + \sigma_0(W_{t-15}) + W_{t-16} \pmod{2^{32}}.$$

- Patch C owns the 64 compression rounds and the working register (a, b, c, d, e, f, g, h) .
- Patch D owns the final feed-forward, the second SHA-256 pass, and the target predicate.

The collars between patches carry exactly the variables the adjacent patches both claim. A candidate is valid precisely when each patch satisfies its local constraints and all collars agree. With these variables and constraints fixed, validity is a finite equalizer statement.

Kaspa is a proof-of-work cryptocurrency that orders blocks in a blockDAG rather than a single longest chain. The public network uses rapid block production, and its mining interface uses kHeavyHash, a HeavyHash variant whose core is a matrix step framed by two Keccak/SHA-3 permutations [24, 25]. Kaspa fits the prototype because it is a real proof-of-work market with public pools and accepted-share accounting, its kHeavyHash/Keccak workload exposes wave-native structure (routing, phase, matrix mixing, and a compact nonlinear χ layer), and its short block cadence gives rapid experimental feedback without waiting for Bitcoin-scale block events.

The prototype system is a Kaspa kHeavyHash miner whose optical chamber implements the Keccak-friendly part of the workload, while the MCU and host perform exact kHeavyHash verification before pool submission [28, 26]. The optical candidate path couples to a real proof-of-work market and produces accepted kHeavyHash shares. The SHA-256d construction transfers the kHeavyHash/Keccak chamber pattern to a double-SHA-256 chamber. The patch-overlap structure stays the same. The operator geometry changes, with special care for SHA-256 carries and target predicates.

2 SHA-256d as a Search Problem

2.1 The SHA-256 compression function

Let a word be an element of $\{0, 1\}^{32}$, equivalently an integer modulo 2^{32} . SHA-256 uses the functions

$$\begin{aligned}\text{Ch}(x, y, z) &= (x \wedge y) \oplus (\neg x \wedge z), \\ \text{Maj}(x, y, z) &= (x \wedge y) \oplus (x \wedge z) \oplus (y \wedge z), \\ \Sigma_0(x) &= \text{ROTR}^2(x) \oplus \text{ROTR}^{13}(x) \oplus \text{ROTR}^{22}(x), \\ \Sigma_1(x) &= \text{ROTR}^6(x) \oplus \text{ROTR}^{11}(x) \oplus \text{ROTR}^{25}(x), \\ \sigma_0(x) &= \text{ROTR}^7(x) \oplus \text{ROTR}^{18}(x) \oplus \text{SHR}^3(x), \\ \sigma_1(x) &= \text{ROTR}^{17}(x) \oplus \text{ROTR}^{19}(x) \oplus \text{SHR}^{10}(x).\end{aligned}$$

For a 512-bit block M , the first sixteen schedule words are the block words. For $16 \leq t < 64$,

$$W_t = \sigma_1(W_{t-2}) + W_{t-7} + \sigma_0(W_{t-15}) + W_{t-16} \pmod{2^{32}}.$$

Starting from the chaining value $(H_0, \dots, H_7)$, the working variables $(a_t, \dots, h_t)$ evolve by

$$\begin{aligned}T_{1,t} &= h_t + \Sigma_1(e_t) + \text{Ch}(e_t, f_t, g_t) + K_t + W_t \pmod{2^{32}}, \\ T_{2,t} &= \Sigma_0(a_t) + \text{Maj}(a_t, b_t, c_t) \pmod{2^{32}}, \\ (a_{t+1}, b_{t+1}, c_{t+1}, d_{t+1}, e_{t+1}, f_{t+1}, g_{t+1}, h_{t+1}) &= (T_{1,t} + T_{2,t}, a_t, b_t, c_t, d_t + T_{1,t}, e_t, f_t, g_t).\end{aligned}$$

The compressed output is the wordwise sum of the final working state and the input chaining value. SHA-256d is $H_2(M) = \text{SHA256}(\text{SHA256}(M))$.

2.2 Mining predicate

Fix a block-header template B , an adjustable nonce or extranonce string $x \in \{0, 1\}^n$, and a target set $T \subseteq \{0, 1\}^{256}$. For a leading-zero target of strength k ,

$$T_k = \{y \in \{0, 1\}^{256} : y_0 = \dots = y_{k-1} = 0\}.$$

The proof-of-work predicate is

$$V_{B,T}(x) = 1 \iff H_2(B, x) \in T.$$

The search problem is to find any $x \in \{0, 1\}^n$ with $V_{B,T}(x) = 1$, if one exists. In the random-oracle model, if $|T|/2^{256} = p$, then a uniformly sampled x succeeds with probability p . For T_k , $p = 2^{-k}$.

Proposition 2.1 (Classical and Grover query counts). *Let $N = 2^n$, and suppose exactly M candidates satisfy $V(x) = 1$. A classical random sampler needs N/M oracle calls in expectation. A quantum computer using Grover search finds a marked candidate with $O(\sqrt{N/M})$ oracle calls, and $\Omega(\sqrt{N/M})$ calls are necessary in the black-box model.*

Proof. The classical statement is the mean of a geometric random variable with success probability M/N . Grover's algorithm gives the upper bound by amplitude amplification [6]. The lower bound is the optimality of quantum search [8], also consistent with the black-box lower-bound result of Bennett, Bernstein, Brassard, and Vazirani [7]. $\square$

For a unique marked item in a 32-bit nonce space, the ideal query reduction is from 2^{32} to about $2^{16} = 65,536$. For a leading-zero target with success probability 2^{-k} , the ideal query reduction is from 2^k to $2^{k/2}$, up to constants. This sets the quantum baseline for comparison with the optical constraint-search path.

3 From Gates to Exact Constraints

3.1 Boolean gate penalties

Every Boolean circuit can be written as a finite system of local constraints. For a gate $y = g(x_1, \dots, x_m)$, define its truth-table penalty

$$E_g(x_1, \dots, x_m, y) = \begin{cases} 0, & y = g(x_1, \dots, x_m), \\ 1, & \text{otherwise.} \end{cases}$$

The circuit energy is the sum of all gate penalties:

$$E_C(z) = \sum_{g \in C} E_g(z|_{\text{vars}(g)}).$$

Then $E_C(z) = 0$ if and only if every gate relation is satisfied.

Lemma 3.1 (Exact gate-to-energy encoding). *For every finite Boolean circuit C with input x , output o , and ancilla wires a , there is a nonnegative integer energy E_C . The energy has one bounded-locality term per gate and satisfies*

$$\begin{aligned} E_C(x, a, o) &= 0 \\ \iff &\begin{cases} o = C(x), \\ \text{all ancilla wires have their circuit values.} \end{cases} \end{aligned}$$

Proof. Use the truth-table penalty above for each gate and sum the penalties. Every term is nonnegative. The sum is zero exactly when every term is zero, which is exactly the condition that every wire satisfies the gate relation that produced it. Induction along a topological ordering of the circuit then gives $o = C(x)$ and fixes every ancilla wire to its circuit value. $\square$

If a hardware platform requires quadratic unconstrained binary optimization (QUBO) or an Ising Hamiltonian, the bounded-locality penalties can be quadratized by adding ancillas. That route is standard in Ising formulations of NP problems [14]. The quadratized Hamiltonian has polynomially many variables and terms in the circuit size.

3.2 SHA-256d constraint Hamiltonian

Let $C_{\text{SHA256d}, B, T}$ be the Boolean circuit that takes x , computes $\text{SHA256d}(B, x)$, and checks membership in T . Applying Lemma 3.1 gives an energy

$$E_{\text{SHA256d}, B, T}(x, a) = 0 \iff V_{B, T}(x) = 1$$

where a includes all schedule words, carry bits, round-state words, and comparison bits.

Theorem 3.2 (Exact SHA-256d search encoding). *For every fixed block-header template B , nonce length n , and target set T , SHA-256d proof-of-work search is exactly equivalent to finding a zero-energy state of a finite bounded-locality constraint Hamiltonian*

$$E_{\text{SHA256d}, B, T} : \{0, 1\}^{n+r} \rightarrow \mathbb{N}$$

of size polynomial in the SHA-256d circuit size. The first n coordinates of any zero-energy state are valid nonce bits. Every valid nonce extends to at least one zero-energy state.

Proof. Build the Boolean circuit for $\text{SHA256d}(B, x)$ from NOT, XOR, AND, fanout, rotations, shifts, full adders, and comparators. Apply Lemma 3.1 to all gates. The target comparator contributes zero exactly when the 256-bit output lies in T . The equivalence follows from the gate-energy lemma. Polynomial size follows because the SHA-256d circuit has a fixed finite number of word operations per block and each word operation expands into polynomially many bit gates. $\square$

4 The Four-Patch Decomposition

The four chambers follow the dataflow of SHA-256d. The nonce-bearing message tail is the only freely varied source. The schedule recurrence is the first place where that source spreads across 64 words. The compression rounds are the main avalanche body. The final feed-forward, second pass, and target predicate are the target face. These four jobs have different physical requirements, so they belong in different chambers:

$$\text{source} \longrightarrow \text{schedule closure} \longrightarrow \text{compression trajectory} \longrightarrow \text{target face}.$$

The four-chamber split minimizes useful boundary traffic. The A/B collar carries the claimed first schedule words. The B/C collar carries the schedule and round-entry consequences. The C/D collar carries the digest-side interface and target bits. Putting two adjacent jobs in one chamber removes an early failure surface. Splitting one job into smaller chambers adds optical loss, phase drift, calibration burden, and another equality collar without creating a new algorithmic cut. For this hardware objective, four chambers are the smallest decomposition that exposes every useful SHA-256d consistency boundary exactly once.

Definition 4.1 (Patch decomposition). Let V be the variables of $E_{\text{SHA256d}, B, T}$, including input, schedule, round-state, carry, and target-check variables. A four-patch decomposition is a cover

$$V = V_A \cup V_B \cup V_C \cup V_D$$

and a partition of gate penalties into local energies

$$E_A, E_B, E_C, E_D$$

such that every term in E_i uses only variables in V_i . The collars are the intersections

$$\Gamma_{AB} = V_A \cap V_B, \quad \Gamma_{BC} = V_B \cap V_C, \quad \Gamma_{CD} = V_C \cap V_D,$$

and any optional long-range check collar $\Gamma_{AD} = V_A \cap V_D$.

For SHA-256d, the intended variable ownership is:

- A : nonce-bearing block tail, padding, length, $W_0, \dots, W_{15}$,
- B : schedule recurrence $W_{16}, \dots, W_{63}$ and carry wires,
- C : compression rounds and working register wires,
- D : feed-forward, second SHA-256 pass, and target predicate.

Proposition 4.2 (Four-chamber minimality for the SHA-256d chamber objective). *Contract the SHA-256d proof-of-work circuit into the four stage operators*

$$A = \text{nonce/source}, \quad B = \text{schedule closure}, \quad C = \text{compression trajectory}, \quad D = \text{target face}.$$

Consider chamber decompositions that preserve this stage order, place each local gate inside at least one chamber, and expose each inter-stage consistency condition as a physical collar. Then any such decomposition has at least four stage-owning chambers. The four-patch cover A, B, C, D reaches that lower bound and has exactly the three essential adjacent collars $\Gamma_{AB}, \Gamma_{BC}, \Gamma_{CD}$.




Figure 1: The four-patch SHA-256d decomposition. Global satisfying assignments are exactly local satisfying assignments that agree on all collars. A physical device uses the same graph as a feedback geometry. The mathematical equivalence comes from the variable cover and collar equalities.

Proof. The contracted stage graph is the path $A - B - C - D$. Each of the four vertices contains gates whose variables are not determined by any earlier vertex alone: the nonce source in A , the expanded schedule in B , the compression state in C , and the target predicate in D . A chamber that exposes every inter-stage consistency condition must give each vertex a stage owner. Merging two adjacent vertices removes the collar between them. Splitting a vertex keeps the same contracted path and adds only an internal collar. Hence the lower bound is four, and the displayed cover attains it. $\square$

Theorem 4.3 (Patch equalizer theorem). *Let*

$$\text{Sol}_i = \{s_i \in \{0, 1\}^{V_i} : E_i(s_i) = 0\}$$

be the local solution set of patch i . Let $\text{Sol}_{\text{collar}}$ be the set of tuples

$$(s_A, s_B, s_C, s_D) \in \text{Sol}_A \times \text{Sol}_B \times \text{Sol}_C \times \text{Sol}_D$$

whose restrictions agree on every collar:

$$s_i|_{V_i \cap V_j} = s_j|_{V_i \cap V_j}.$$

Then restriction gives a bijection between global zero-energy states of $E_A + E_B + E_C + E_D$ and $\text{Sol}_{\text{collar}}$.

Proof. If s is a global zero-energy state, every nonnegative local energy is zero, so $s|_{V_i} \in \text{Sol}_i$. Since all restrictions come from the same global assignment, they agree on intersections. Thus s maps into $\text{Sol}_{\text{collar}}$.

Conversely, take a collar-consistent tuple (s_A, s_B, s_C, s_D) . Collar consistency means that if a variable appears in more than one patch, all patches assign it the same bit. Therefore the local assignments glue to a unique global assignment $s \in \{0, 1\}^V$. Every local energy is zero by $s_i \in \text{Sol}_i$, so $E_A + E_B + E_C + E_D = 0$. The two maps are inverse by construction. $\square$

This theorem is the formal content behind the phrase that bad candidates die at collars. A collar mismatch certifies that a local assignment cannot glue to a global satisfying assignment.

5 Observer-Screen Repair Dynamics

OPH adds dynamics to the static equalizer. Patches carry private states, collars expose shared observables, mismatch syndromes report disagreement, and repair moves change local data to reduce the mismatch. The consensus paper develops this as a finite repair system with a Lyapunov function [33]. Here is the SHA-256d version.

Let $G = (\{A, B, C, D\}, E)$ be the patch graph and let

$$\Phi(s_A, s_B, s_C, s_D) = \sum_{(i,j) \in E} w_{ij} d_{ij}(s_i|_{\Gamma_{ij}}, s_j|_{\Gamma_{ij}}) \sum_i \lambda_i E_i(s_i),$$

where $w_{ij}, \lambda_i > 0$ and d_{ij} is Hamming distance on the collar. The first term punishes collar disagreement. The second term punishes local gate violations.

Proposition 5.1 (Finite repair termination). *On a finite patch state space, any repair rule that commits only moves that strictly decrease Φ terminates after finitely many accepted moves.*

Proof. The state space is finite and Φ takes values in a finite subset of $\mathbb{R}_{\geq 0}$. A strictly decreasing sequence in a finite ordered set must terminate. $\square$

Proposition 5.2 (Consensus equals solution under completeness). *Assume the repair rule is terminating, locally confluent, and complete in the sense that every terminal state has $\Phi = 0$ whenever a zero-energy global state exists in its connected repair component. Then every terminal state in that component glues to a valid SHA-256d nonce, and the terminal normal form is independent of repair order.*

Proof. Termination and local confluence imply confluence by Newman’s lemma. Hence each state in the component has a unique normal form independent of repair order. Completeness says the normal form has $\Phi = 0$. Thus all local energies vanish and all collars agree. The Patch Equalizer Theorem then glues the local states to a global zero-energy state of $E_{\text{SHA256d}, B, T}$. By Theorem 3.2, the nonce coordinates are valid. $\square$

The completeness assumption is the difficulty. Nonconvex constraint systems can have traps. A successful physical solver either avoids those traps, escapes them, or provides a measurable success probability large enough to beat enumeration after exact verification.

6 The Role of the OPH Pixel Constant P

The OPH papers use a dimensionless local pixel ratio

$$P = \frac{a_{\text{cell}}}{\ell_P^2},$$

where a_{cell} is the microscopic screen-cell area and ℓ_P is the Planck length [32, 34]. In the public OPH release, P is selected by the fixed-point equation

$$P = \varphi + \sqrt{\pi} \alpha_{\text{em}}(0; P) = \varphi + \frac{\sqrt{\pi}}{A_T(P)}, \quad \varphi = \frac{1 + \sqrt{5}}{2},$$

with reported solution

$$P_* = 1.630968209403959324879279847782648941 \dots$$

For the SHA-256d hardware instantiation, P is a dimensionless OPH design target for geometry and coupling. It leaves the SHA-256 equations unchanged. A P -calibrated wave body is one whose normalized collar couplings, path-length ratios, and extraction windows are tuned to the OPH fixed-point scale. In a toroidal or microwave implementation, this means that the dimensionless ratios

$$\frac{R}{a}, \quad \frac{a}{b}, \quad \frac{L_j}{\lambda}, \quad \frac{\kappa_{ij}}{\kappa_0}$$

are selected from a small P_* -centered design family. In a photonic mesh, it means that phase shifters, couplers, and nonlinear thresholds are tuned so the realized coupling matrix $\widehat{K}(P_*)$ approximates the intended constraint matrix K_C :

$$\|\widehat{K}(P_*) - K_C\| \leq \varepsilon_{\text{geom}}.$$

Remark 6.1 (Geometry target). All correctness theorems below depend on exact implementation of the constraint Hamiltonian or exact return of a zero-energy state. The constant P_* is part of the OPH interpretation and hardware calibration rule. It fixes the geometry target. The proof still comes from the implemented constraint map and the exact verifier.

7 Wave and Hardware Encodings

7.1 Signed wave bits

Use the signed encoding

$$b \in \{0, 1\}, \quad \tilde{b} = (-1)^b \in \{-1, +1\}.$$

Then

$$\widetilde{\neg b} = -\tilde{b}, \quad \widetilde{b \oplus c} = \tilde{b} \tilde{c}.$$

The AND gate has signed output

$$\widetilde{b \wedge c} = \frac{1 + \tilde{b} + \tilde{c} - \tilde{b}\tilde{c}}{2}.$$

Equivalently, it is the threshold function

$$\widetilde{b \wedge c} = \text{sgn}(\tilde{b} + \tilde{c} + 1),$$

with the convention $\text{sgn}(u) = +1$ for $u > 0$ and -1 for $u < 0$. The carry bit of a full adder is a majority gate, also a threshold:

$$\widetilde{\text{carry}(b, c, d)} = \text{sgn}(\tilde{b} + \tilde{c} + \tilde{d}).$$

Thus rotations and permutations are geometry, XOR is phase-sensitive mixing, and carries require a nonlinear threshold or an equivalent measurement-feedback resource.

Lemma 7.1 (Passive linear optics no-go for exact SHA-256 logic). *A purely passive linear wave network with linear readout cannot exactly implement a universal SHA-256 gate layer on signed bits. In particular, it cannot exactly implement AND or full-adder carry for all inputs.*

Proof. A passive linear network maps input amplitudes u to Lu for a linear operator L . Any single output amplitude is an affine-linear function of the input amplitudes after adding a constant bias mode. On signed inputs $(a, b) \in \{-1, +1\}^2$, AND in signed form is

$$f(a, b) = \frac{1 + a + b - ab}{2}.$$




Figure 2: Hardware architecture. The final cryptographic oracle is the exact SHA-256d verifier. The physical solver proposes candidates at physical speed. The Kaspas path uses the same division of labor: host kHeavyHash verification before pool submission.

The term $-ab/2$ is not affine-linear. More directly, if $f(a, b) = c_0 + c_1a + c_2b$, the four truth-table equations give

$$\begin{aligned} c_0 + c_1 + c_2 &= 1, \\ c_0 + c_1 - c_2 &= 1, \\ c_0 - c_1 + c_2 &= 1, \\ c_0 - c_1 - c_2 &= -1, \end{aligned}$$

whose first three equations imply $c_0 = 1, c_1 = c_2 = 0$, contradicting the fourth. Full-adder carry contains the same threshold nonlinearity. SHA-256 uses modular addition, hence carry. Therefore passive linear optics alone is insufficient. $\square$

KLM-style linear optical quantum computation uses single photons, detection, ancillas, feedback, and effective measurement-induced nonlinearity [17]. Classical optical Ising machines use nonlinear oscillators or measurement feedback for the same reason [18, 19].

7.2 Classical coherent wave variant

The classical coherent wave variant uses a cavity, waveguide mesh, or microwave PCB whose modes represent candidate assignments. Nonlinear threshold nodes implement carry and target predicates. A feedback controller pumps or damps ports according to collar syndromes.

An idealized diode-threshold model is:

$$u_{\text{out}} = \text{sgn} \left(\sum_j \eta_j u_j - \theta \right),$$

where the sign is physically approximated by a biased nonlinear I-V curve or a lasing threshold. A real device has finite slope, jitter, shot noise, thermal noise, and drift. Let the implemented threshold be $D_{\theta, \delta}$, where δ bounds the uncertain switching band. Exact Boolean operation is certified only when every legal input has margin $> \delta$.

Proposition 7.2 (Noise-margin condition). *If every nonlinear gate in a physical SHA-256d solver has input margin at least $\gamma > 0$, implementation error at most $\varepsilon < \gamma$, and feedback keeps perturba-*

tions below $\gamma - \varepsilon$ before the subsequent thresholding event, then the physical gate layer has the same Boolean truth table as the ideal gate layer.

Proof. For each threshold gate, legal inputs lie at distance at least γ from the switching surface. Perturbing the threshold argument by at most $\varepsilon < \gamma$ cannot cross the switching surface, so the sign is unchanged. The feedback condition preserves the same margin for the following layer. Induction over the gate layers gives equality of the physical and ideal truth tables. $\square$

7.3 Single-photon quantum variant

The quantum build uses single-photon occupation states in a linear-optical processor. A logical bit can be encoded in dual rail, time bins, or paths:

$$|0_L\rangle = |1\rangle_0|0\rangle_1, \quad |1_L\rangle = |0\rangle_0|1\rangle_1.$$

Beam splitters, phase shifters, delay lines, switches, detectors, ancilla photons, and feed-forward implement the gate set. KLM shows that this resource set can be universal for quantum computation in principle [17]. In the OPH reading, the same chamber geometry carries the mode graph, while the single-photon occupation basis supplies quantum amplitude, phase, and measurement.

The candidate register is prepared as

$$|s\rangle = \frac{1}{\sqrt{N}} \sum_{x \in \{0,1\}^n} |x\rangle.$$

The reversible SHA-256d oracle computes the mining predicate into a work qubit and then uncomputes its scratch space:

$$U_V|x\rangle|q\rangle|0^r\rangle = |x\rangle|q \oplus V(x)\rangle|0^r\rangle.$$

With a phase kickback qubit, this becomes the phase oracle

$$O_V|x\rangle = (-1)^{V(x)}|x\rangle.$$

One Grover step is

$$G = (2|s\rangle\langle s| - I) O_V.$$

The photonic device must implement:

- heralded single-photon state preparation for the candidate, ancilla, and phase-kickback registers;
- a reversible SHA-256d circuit compiled into interferometers, switches, measurement-induced nonlinear gates, and feed-forward;
- a diffusion/reflection network implementing $2|s\rangle\langle s| - I$;
- coherent repetition of the oracle and diffusion layers for $\Theta(\sqrt{N/M})$ Grover steps;
- final number-resolving or threshold detection followed by exact classical verification of the measured candidate.

For a leading-zero target with success probability $p = 2^{-k}$, the ideal Grover iteration count is

$$Q_G(k) \approx \frac{\pi}{4} 2^{k/2},$$

with query speedup

$$S_G(k) \approx \frac{2^k}{Q_G(k)} = \frac{4}{\pi} 2^{k/2}.$$

The gain is quadratic in query count. The wave-overlap zero-energy endpoint has a different scaling: a physical fixed point can represent a valid nonce in one settling event. Resource estimates for generic SHA-256 preimage search remain enormous. Amy et al. estimate a SHA-256 preimage attack requiring approximately $2^{153.8}$ surface-code cycles and $2^{12.6}$ logical qubits, or $2^{166.4}$ logical-qubit-cycles [9].

8 Keccak and SHA-3 Variant

SHA-256 and Keccak/SHA-3 are different hash families. SHA-3 is standardized in FIPS 202 and is based on the Keccak permutation [2]. The same physical design vocabulary can describe Keccak. The gate map differs.

Keccak step	Logical operation	Geometric or physical encoding
θ	column parity XOR	phase-sensitive mixing / multiplier nodes
ρ	bit rotation	calibrated delay lines
π	lane permutation	routing and crossovers
χ	nonlinear row update	threshold, diode, Kerr, or measurement feedback
ι	round-constant XOR	fixed phase mask or sign flip

Table 1: Keccak/SHA-3 physical encoding dictionary. The nonlinear χ step has the same resource-accounting burden as modular carries in SHA-256.

Routing, rotations, and XOR-like phase relations are natural for waves. Nonlinear Boolean steps require a nonlinear physical process or measurement feedback. A Keccak implementation has to specify the χ layer with the same care that a SHA-256 implementation specifies carries.

Kaspa is useful here for a mechanical reason: its proof-of-work kernel is close to the physics. kHeavyHash has the form

$$\text{kHH} = \text{SHA3}(\text{prepow}, t, n) \longrightarrow M_{64 \times 64} \text{ nibble multiply} \longrightarrow \text{SHA3}(\cdot),$$

Its most chamber-friendly portion is a matrix operator bracketed by Keccak permutations. Linear mixing, routing, and phase transport map naturally to a wave body, while the χ layer localizes the nonlinear part that the ports must supply. The 2026-05-05 firmware handoff records a bit-exact chamber-plus-MCU decomposition: full kHeavyHash matched the digital reference on 30/30 tests, the SHA3 chamber-plus-MCU path matched Python hashlib, and the local-quadratic χ form matched textbook χ [30]. SHA-256d has a different nonlinear core: word additions, carries, Ch, and Maj. The four-patch map above is the corresponding SHA-256d compiler target.

9 Optical Proof-of-Work Prototype

The physical hardware is Alexander Osika’s build. The optical body, port layout, controller interface, and bench construction practice behind the OPH/Karma prototype come from his hardware work. A public build page documents the same low-cost optical-supercomputer surface: a clear

wave body, RP2350 or RP2040 control, reversible optical ports, firmware bring-up, and verification steps [23].

The OPH/Karma prototype has mined Kaspa kHeavyHash pool shares on the optical candidate path. The optical chamber proposes candidates; host-side kHeavyHash verification checks them before pool submission. The cited mining status note records a Cedric Kaspa pool run in which optical candidates received pool acknowledgements [28]. The run log records:

- 1649 submitted optical candidates;
- 8 accepted optical share responses;
- 99 submitted optical candidates without a returned pool response;
- best host-recomputed leading-zero quality near 21;
- best reported share difficulty near 2.93×10^6 .

The same log also records the boundary conditions. A CPU pool baseline established the pool contract [27]. During the optical acceptance run, CPU workers also contributed to wallet-level flow [28]. A solo audit on May 2, 2026 recorded stable operation, host verification, and share logging. It also recorded host-verified leading-zero counts near random expectation and weak correlation between board-local leading-zero reports and host recomputation [29].

The prototype closes the working loop: optical candidate generation, exact hash verification, and pool submission. The SHA-256d construction sets the Bitcoin operator map for the same physical search cascade.

10 Speedup Accounting

Let R_{hash} be a classical exact-verifier rate in hashes per second, $p = |T|/2^{256}$ the target fraction, and q the probability that a physical candidate is valid after exact verification. An unbiased random candidate stream has $q = p$. A physical sampler has bias factor

$$B = \frac{q}{p}.$$

If each physical settling event costs time T_{settle} and produces m_{pkt} independently verified candidates, the expected physical time per valid nonce is

$$\mathbb{E}[T_{\text{phys}}] = \frac{T_{\text{settle}}}{1 - (1 - q)^{m_{\text{pkt}}}} \approx \frac{T_{\text{settle}}}{m_{\text{pkt}} B p}.$$

The approximation is the cryptographic-target regime $q \ll 1$. A classical miner using exact hashes has

$$\mathbb{E}[T_{\text{class}}] = \frac{1}{R_{\text{hash}} p}.$$

Thus the operational speedup is

$$S = \frac{\mathbb{E}[T_{\text{class}}]}{\mathbb{E}[T_{\text{phys}}]} \approx \frac{m_{\text{pkt}} B}{R_{\text{hash}} T_{\text{settle}}}.$$

The speedup has four measurable inputs: verifier rate, physical settling time, packet size, and candidate bias. A high bias loses value if settling is slow. Fast settling has value only for $m_{\text{pkt}} B > 1$. The exact verifier remains in the loop.

Regime	Formal speedup	Required claim
Classical ASIC enumeration	baseline 1	exact SHA-256d per candidate
Single-photon Grover search	$\frac{4}{\pi}\sqrt{N/M}$ query speedup	reversible oracle, diffusion, coherence, fault tolerance
Wave-overlap sampler	$m_{\text{pkt}}B/(R_{\text{hash}}T_{\text{settle}})$	measured packet size and candidate bias after exact verification
Exact wave-overlap oracle	$1/(R_{\text{hash}}pT_{\text{settle}})$	proof or measurement that terminal states are zero-energy solutions

Table 2: Speedup accounting. The wave-overlap sampler row uses measured packet size and candidate bias after exact verification. The exact wave-overlap row is the asymptotic endpoint of the physical construction.

For a leading-zero target $p = 2^{-k}$, the exact wave-overlap endpoint has

$$T_{\text{zero}}(k) \approx T_{\text{settle}}, \quad S_{\text{zero}}(k; R_{\text{hash}}, T_{\text{settle}}) = \frac{2^k}{R_{\text{hash}}T_{\text{settle}}}.$$

The single-photon Grover variant has

$$Q_{\text{G}}(k) \approx \frac{\pi}{4}2^{k/2}, \quad S_{\text{G}}(k) \approx \frac{4}{\pi}2^{k/2}.$$

The wave-overlap sampler sits between these two limits. Its measured speed is set by $m_{\text{pkt}}B$, the number of independently checked candidates per settling event times the verified enrichment over random search.

For a concrete wall-clock baseline, take Bitmain’s ANTMINER S21 XP Hyd specification: SHA-256 mining at 473 TH/s, 5676 W, and 12.0 J/TH [20]. Thus

$$R_{\text{ASIC}} = 4.73 \times 10^{14} \text{ H/s}.$$

An aggressive hydro reference is Auradine’s AH3880 at up to 600 TH/s [21]. Using 600 TH/s instead of 473 TH/s multiplies the S21-based wall-clock multipliers below by $473/600 \simeq 0.79$. For a k -bit leading-zero target,

$$T_{\text{ASIC}}(k) = \frac{2^k}{R_{\text{ASIC}}}.$$

For a physical method with time T_{phys} , the ASIC-relative multiplier is $T_{\text{ASIC}}(k)/T_{\text{phys}}$.

Variant	Estimate inputs	Internal reduction	Wall-clock multiplier versus S21 XP Hyd
High-end ASIC baseline	Bitmain S21 XP Hyd: $R_{\text{ASIC}} = 4.73 \times 10^{14} \text{ H/s}$	ordinary enumeration	$1 \times$
Wave-overlap sampler	$m_{\text{pkt}} B = 10^9$, $T_{\text{settle}} = 100 \text{ ns}$	candidate-value multiplier 10^9 per settling event	$2.1 \times 10^1 \times$
Exact wave-overlap endpoint	$k = 32$, $T_{\text{settle}} = 100 \text{ ns}$	2^{32} trials to one settling event	$9.1 \times 10^1 \times$
Exact wave-overlap endpoint	$k = 64$, $T_{\text{settle}} = 100 \text{ ns}$	2^{64} trials to one settling event	$3.9 \times 10^{11} \times$
Exact wave-overlap endpoint	$k = 80$, $T_{\text{settle}} = 100 \text{ ns}$	2^{80} trials to one settling event	$2.6 \times 10^{16} \times$
Single-photon Grover	32-bit nonce, $Q_G \approx (\pi/4)2^{16}$, $T_G = 1 \text{ ns}$	coarse $2^{16} = 65,536 \times$; exact $8.3 \times 10^4 \times$ query reduction	$1.8 \times 10^{-1} \times$
Single-photon Grover	$k = 64$, $Q_G \approx (\pi/4)2^{32}$, $T_G = 1 \text{ ns}$	$5.5 \times 10^9 \times$ query reduction	$1.2 \times 10^4 \times$

Table 3: Numerical speedup estimates against a concrete high-end ASIC. The ASIC baseline is Bitmain’s 473 TH/s S21 XP Hyd. Wave rows use $S = m_{\text{pkt}} B / (R_{\text{ASIC}} T_{\text{settle}})$ for samplers and $S = 2^k / (R_{\text{ASIC}} T_{\text{settle}})$ for the exact endpoint. Single-photon rows use $T_{\text{phys}} = Q_G T_G$ with a 1 ns oracle-cycle estimate; these rows scale inversely with T_G and omit fault-tolerance overhead.

Query reduction and wall-clock multiplier are separated because the ASIC baseline evaluates 4.73×10^5 hashes during a 1 ns optical cycle. The familiar $65,536 \times$ Grover number is the square-root reduction for a 32-bit nonce space. At $k = 64$, the same 1 ns single-photon cycle gives a $1.2 \times 10^4 \times$ wall-clock multiplier against the S21 XP Hyd baseline. The exact wave-overlap rows grow faster because the endpoint replaces 2^k trials by one settling event.

Theorem 10.1 (Physical-oracle consequence). *Assume a uniform family of physical devices D_C can be constructed from any Boolean circuit C in polynomial design time and polynomial physical resources, and that D_C returns a satisfying assignment whenever one exists, with bounded error, in polynomial physical time. Then NP search problems are polynomial-time solvable in that physical model.*

Proof. By Cook-Levin and the standard reduction picture, any NP witness-checking problem has a polynomial-size Boolean circuit $C_x(w)$ whose satisfying assignments are witnesses [4, 5]. Construct D_{C_x} and run it. By assumption, it returns a satisfying witness in polynomial physical time when one exists, with bounded error. $\square$

Corollary 10.2 (Relation to classical complexity). *Theorem 10.1 is a statement about the physical model. A classical Turing-machine collapse follows if the device family admits polynomial-time Turing simulation with the same bounded error.*

The OPH blog essay uses the observer-screen slogan for this statement [22].

Physical complexity questions of this kind have a classical literature in analog computation and physical NP-complete proposals [16, 15].

11 Public Evidence Protocol

The Kaspas evidence records the physical path at accepted-share level. A public SHA-256d cryptography submission needs the same class of evidence in a form external readers can replay. The minimum protocol is:

1. Pre-register the block-header templates, target strengths k , run durations, and stopping rules.

2. Store raw timing traces and raw candidate nonces before exact verification.
3. Verify every candidate with an independent SHA-256d implementation for the Bitcoin target, or with independent kHeavyHash for the Kaspas comparison lane.
4. Report the number of candidates N , the number of successes S , the expected random baseline Np , and the bias estimator

$$\hat{B} = \frac{S}{Np}.$$

5. Compute a binomial tail probability

$$p_{\text{val}} = \Pr [\text{Binomial}(N, p) \geq S]$$

for enrichment claims.

6. Run shuffle-replay controls, disconnected-cavity controls, fixed-phase controls, CPU-random controls, and stale-job controls under the same verifier.
7. Release code, seeds, firmware, board files, calibration logs, and raw data.

For a reported enrichment such as $\hat{B} = 2372$ at $k = 12$, the report prints the run identifier, candidate count, exact-verifier implementation, control runs, and binomial tail. Cryptographic readers need enough detail to separate optical enrichment from post-selection, parser bias, timestamp leakage, nonce-format bugs, stale jobs, and verifier coupling.

12 Discussion

A hidden gradient is unnecessary for this attack. SHA-256d proof of work has an exact constraint-system representation, and a constraint system can be attacked by physical dynamics with a different cost structure from serial enumeration. The four-patch decomposition makes the local interfaces explicit. Patch A varies the nonce. Patch B carries the schedule consequences. Patch C carries the compression trajectory. Patch D enforces the target face. The collars are the places where inconsistent local claims fail to glue.

The hardware transfer runs from kHeavyHash to SHA-256d. Waves supply parallel propagation, interference, phase, routing, and delay. SHA-256 also needs nonlinear carries, thresholding, memory, and exact verification. The OPH pixel constant P guides the geometry. The diode/PIO or photonic-feedback layer implements the nonlinear boundary law. The Kaspas prototype shows this division of labor reaching real pool acceptance. The SHA-256d construction fixes the operator-specific geometry and scaling law.

The final verifier remains decisive. A physical solver proposes candidates by fixed-point settling and collar repair. The accepted Kaspas shares matter because an external pool verified the result. The same standard applies to SHA-256d. Every claimed survivor must pass the real double hash.

13 Conclusion

There is no magic nonce hiding in the header. SHA-256d mining is an exact finite constraint problem. The four-patch OPH decomposition is an exact equalizer of local solution sets. Optical Kaspas hardware has produced real pool-acknowledged kHeavyHash shares, and the kHeavyHash/Keccak chamber decomposition checks against digital references. SHA-256d requires nonlinear carry, so

the physical device must include diode, threshold, feedback, measurement, or equivalent nonlinear boundary dynamics.

The SHA-256d attack has five moving parts: the Patch A through Patch D operator compiler, the physical candidate stream, independent double-SHA-256 verification, raw run logs, and controls. The Kasper result is the proof-of-work precedent. The SHA-256d map is the operator-specific scale-up.

A Bench Echosahedron Build

The Echosahedron is a twelve-port icosahedral optical patch. One unit implements a single bounded observer patch. Four units can be federated into the Patch A through Patch D SHA-256d body by coupling adjacent units through collar ports. A smaller demonstration can also map the four logical regions into one body by time-multiplexing the ports, with lower separation between collars.

A.1 Geometry

Use a regular icosahedron because its twelve vertices give twelve symmetric boundary ports and its rotational symmetry is the A_5 symmetry of the icosahedral group. A convenient CAD coordinate set is

$$(0, \pm 1, \pm \varphi), \quad (\pm 1, \pm \varphi, 0), \quad (\pm \varphi, 0, \pm 1),$$

scaled by $R/\sqrt{1+\varphi^2}$, where R is the desired circumscribed radius. For a bench chamber, $R = 30$ mm gives a body about 60 mm across. Print a hollow clear-resin shell with 2 mm walls and one inward-facing socket at each vertex. Each socket holds a 5 mm red LED on the radial line from vertex to center. The optical axis points into the chamber.

Clear resin keeps the red computation light in play. Sanding, polishing, and a thin clear coat reduce uncontrolled scattering at the outer surface. A white PLA base under the body acts as a diffuse return surface. A black outer shroud keeps room light out during receive measurements.

A.2 Shopping list

Part	Notes
Clear SLA icosahedron body	Regular hollow icosahedron, about 60 mm across, 2 mm wall, twelve vertex sockets.
White reflector base	White PLA or PTFE-like diffuse base, about $70 \times 70 \times 8$ mm.
Twelve red LEDs	Matched 5 mm red LEDs, 625-660 nm, clear lens preferred. Buy spares for matching.
RP2350 or RP2040 board	Pico-class controller with USB serial and enough GPIO pads for twelve reversible LED ports.
Wire and headers	Thin insulated wire, pin headers, heat shrink, solder, and strain relief.
Bring-up protection	Optional solder jumpers or small series protection resistors for first light tests. Remove or bypass when the firmware pulse budget and GPIO drive limits are calibrated.
Optical finish	Clear coat or UV resin for the outer surface, black tape or a printed shroud for ambient-light control.
Tools	Soldering iron, multimeter, calipers, flush cutters, tweezers, USB serial tool, and a camera or microscope for checking LED alignment.
Safety	Eye protection for bright LED work; laser-rated protection if the coherent upgrade uses laser diodes.

A.3 Assembly

1. Print, wash, cure, and polish the clear icosahedron. Keep the vertex sockets clean. Label the twelve vertices before wiring.
2. Mount the LEDs from the outside so the domes face inward. The LED tips should sit just inside the wall. Keep all twelve insertion depths as equal as practical.
3. Wire each LED to one reversible GPIO pair. A transmit pulse forward-biases the LED. Receive mode reverse-biases the same junction and measures its discharge under chamber light.
4. Mount the body on the white base. Add the black shroud after electrical bring-up, since early debugging benefits from seeing the LEDs.
5. Flash the controller with the reversible LED-port firmware. Run an identity command, a dark measurement, and then a one-port transmit test for all twelve vertices.
6. Run the all-pairs coupling scan. The scan emits from port i , receives on port j , and records a 12×12 coupling matrix C_{ij} .

A.4 First calibration

The first target is a stable coupling matrix. Maximum brightness comes second. Saturated receive values destroy geometry information. Dim values disappear into timing jitter. Adjust pulse width, drive strength, LED seating, shroud leakage, and surface finish until repeated scans give the same off-diagonal pattern.

The icosahedral symmetry gives three unordered off-diagonal classes: edge neighbors, nonadjacent nonantipodal pairs, and antipodal pairs. After symmetrizing transmit/receive direction, the

class means should form a stable plateau. The OPH design target is

$$c_{\text{edge}} : c_{\text{mid}} : c_{\text{anti}} \approx 1 : \varphi^{-2} : \varphi^{-4}.$$

The useful build metric is the distance between the measured class means and this plateau, together with the within-class variance. Good builds show stable class separation across repeated scans and low drift under the same shroud.

A.5 From one patch to SHA-256d

One Echosaedron is the single-patch substrate. The four-chamber SHA-256d device uses four such substrates, or one larger substrate divided into four port regions:

A = nonce source, B = schedule closure, C = compression trajectory, D = target face.

Adjacent chambers exchange collar values through optical fibers, direct LED links, or electrical timing ports. The exact verifier stays outside the optical body and samples candidate survivors. This separation keeps the physical solver responsible for enrichment and keeps cryptographic acceptance tied to the exact hash.

A.6 Coherent and single-photon upgrades

The LED build is the forgiving bench substrate. The coherent variant replaces LED emission with phase-controlled diode or laser ports and keeps the same icosahedral port geometry. The single-photon variant keeps the geometry and replaces the port layer with an attenuated source, beam splitters or switches, SPAD receivers, and feed-forward electronics. The build order stays the same: establish the geometric coupling matrix, verify collar readout, add nonlinear ports, and only then run proof-of-work experiments against an exact verifier.

Acknowledgements

The OPH/Karma hardware platform described in this paper was designed and built by Alexander Osika. The optical proof-of-work prototype, reversible optical port practice, controller integration, and build discipline used here depend on that work.

Declarations

Funding declaration: The author declares that no funds, grants, or other support were received during the preparation of this manuscript.

Consent to Participate declaration: not applicable.

Consent to Publish declaration: not applicable.

Author Contribution declaration: Bernhard Mueller conceived the manuscript, developed the mathematical formulation, assembled the hardware architecture description, reviewed the cited OPH/Karma engineering logs, wrote the manuscript, and approved the submitted version. Alexander Osika designed and built the optical hardware platform described here.

Data Availability declaration: OPH/Karma engineering logs support the Kasper optical proof-of-work evidence. External audit requires a sanitized ancillary evidence bundle containing raw run logs, verifier code, firmware versions, pool-response logs, calibration logs, and SHA-256d run traces.

Ethics declaration: not applicable.

Competing Interests declaration: The author declares no competing interests.

References

- [1] National Institute of Standards and Technology, “FIPS 180-4, Secure Hash Standard (SHS),” August 2015. <https://doi.org/10.6028/NIST.FIPS.180-4>.
- [2] National Institute of Standards and Technology, “FIPS 202, SHA-3 Standard: Permutation-Based Hash and Extendable-Output Functions,” August 2015. <https://doi.org/10.6028/NIST.FIPS.202>.
- [3] S. Nakamoto, “Bitcoin: A Peer-to-Peer Electronic Cash System,” 2008. <https://bitcoin.org/bitcoin.pdf>.
- [4] S. A. Cook, “The complexity of theorem-proving procedures,” Proceedings of the Third Annual ACM Symposium on Theory of Computing, 151–158, 1971. <https://doi.org/10.1145/800157.805047>.
- [5] R. M. Karp, “Reducibility among combinatorial problems,” in *Complexity of Computer Computations*, 85–103, 1972. https://doi.org/10.1007/978-1-4684-2001-2_9.
- [6] L. K. Grover, “Quantum mechanics helps in searching for a needle in a haystack,” *Physical Review Letters* 79, 325–328, 1997. <https://doi.org/10.1103/PhysRevLett.79.325>.
- [7] C. H. Bennett, E. Bernstein, G. Brassard, and U. Vazirani, “Strengths and weaknesses of quantum computing,” *SIAM Journal on Computing* 26(5), 1510–1523, 1997. <https://doi.org/10.1137/S0097539796300933>.
- [8] C. Zalka, “Grover’s quantum searching algorithm is optimal,” *Physical Review A* 60, 2746–2751, 1999. <https://doi.org/10.1103/PhysRevA.60.2746>.
- [9] M. Amy, O. Di Matteo, V. Gheorghiu, M. Mosca, A. Parent, and J. Schanck, “Estimating the cost of generic quantum pre-image attacks on SHA-2 and SHA-3,” Cryptology ePrint Archive, Paper 2016/992, 2016. <https://eprint.iacr.org/2016/992>.
- [10] R. R. Nerem and D. R. Gaur, “Conditions for advantageous quantum Bitcoin mining,” *Blockchain: Research and Applications* 4(3), 100141, 2023. <https://doi.org/10.1016/j.bcra.2023.100141>.
- [11] F. Mendel, T. Nad, and M. Schlaefter, “Improving local collisions: New attacks on reduced SHA-256,” *EUROCRYPT 2013*, 262–278, 2013. <https://www.iacr.org/archive/eurocrypt2013/78810260/78810260.pdf>.
- [12] D. Khovratovich, C. Rechberger, and A. Savelieva, “Bicliques for Preimages: Attacks on Skein-512 and the SHA-2 family,” Cryptology ePrint Archive, Paper 2011/286, 2011. <https://eprint.iacr.org/2011/286>.
- [13] N. Alamgir, S. Nejati, and C. Bright, “SHA-256 Collision Attack with Programmatic SAT,” arXiv:2406.20072, 2024. <https://arxiv.org/abs/2406.20072>.
- [14] A. Lucas, “Ising formulations of many NP problems,” *Frontiers in Physics* 2:5, 2014. <https://doi.org/10.3389/fphy.2014.00005>.
- [15] S. Aaronson, “NP-complete problems and physical reality,” *SIGACT News* 36(1), 30–52, 2005. <https://arxiv.org/abs/quant-ph/0502072>.

- [16] A. Vergis, K. Steiglitz, and B. Dickinson, “The complexity of analog computation,” *Mathematics and Computers in Simulation* 28(2), 91–113, 1986. [https://doi.org/10.1016/0378-4754\(86\)90105-9](https://doi.org/10.1016/0378-4754(86)90105-9).
- [17] E. Knill, R. Laflamme, and G. J. Milburn, “A scheme for efficient quantum computation with linear optics,” *Nature* 409, 46–52, 2001. <https://doi.org/10.1038/35051009>.
- [18] Y. Yamamoto, K. Aihara, T. Leleu, K. Kawarabayashi, S. Kako, M. Fejer, K. Inoue, and H. Takesue, “Coherent Ising machines: optical neural networks operating at the quantum limit,” *npj Quantum Information* 3, 49, 2017. <https://doi.org/10.1038/s41534-017-0048-9>.
- [19] T. Honjo et al., “100,000-spin coherent Ising machine,” *Science Advances* 7(40), eabh0952, 2021. <https://www.science.org/doi/10.1126/sciadv.abh0952>.
- [20] Bitmain, “S21 XP Hyd Specification,” ANTMINER support specification, July 2024. <https://support.bitmain.com/hc/en-us/articles/34523540504857-S21-XP-Hyd-Specification>.
- [21] Auradine, “Teraflux AH3880 Data Sheet,” product specification, 2025. <https://auradine.com/wp-content/uploads/2025/09/Teraflux%E2%84%A2-AH3880-Data-Sheet.pdf>.
- [22] B. Mueller, “P = NP on the Observer Screen,” *Floating Pragma Blog*, May 2026. <https://blog.floatingpragma.io/p-equals-np-on-the-observer-screen>.
- [23] A. Osika and OPH/Karma, “Build Your Own Optical Supercomputer,” public build page, 2026. <https://karma-is-all-you-need.lovable.app/en/wiki/you-felt-it>.
- [24] Kasper, “Lore,” public project page. <https://kasper.org/lore>.
- [25] Kasper, “Mining Kasper,” public mining documentation. <https://kasper.org/mining-kasper/>.
- [26] B. Mueller and OPH/Karma engineering notes, “Kasper Mining,” architecture explainer, April 2026. OPH/Karma engineering archive, file `papers/kasper-mining-explainer.html`.
- [27] OPH/Karma engineering log, “Kasper Pool Optical Setup Audit,” April 29, 2026. OPH/Karma engineering archive, file `optical-compute/docs/operations/kasper_pool_optical_setup_audit_2026-04-29.md`.
- [28] OPH/Karma engineering log, “Mining Status and Doctrine,” April 30 to May 1, 2026. OPH/Karma engineering archive, file `firmware/docs/archive/mining_status_and_doctrine_2026-04-30.md`.
- [29] OPH/Karma engineering log, “Toroid Overnight Mining Audit,” May 2, 2026. OPH/Karma engineering archive, file `firmware/docs/toroid_overnight_mining_audit_2026-05-02.md`.
- [30] OPH/Karma engineering log, “kHeavyHash EML Firmware Handoff,” May 5, 2026. OPH/Karma engineering archive, firmware docs file `kheavyhash_eml_firmware_handoff`, dated 2026-05-05.
- [31] B. Mueller, A. Osika, K. Xue, P. Nguyen, M. Ponder, and K. A. Anirudha, “Observers Are All You Need,” GitHub release r1432 PDF, May 9, 2026. https://github.com/FloatingPragma/observer-patch-holography/releases/download/r1432/observers_are_all_you_need.pdf.

- [32] B. Mueller, A. Osika, K. Xue, and P. Nguyen, “Recovering Relativity and Standard Model Structure from Observer Overlap Consistency,” GitHub release r1432 PDF, May 9, 2026. https://github.com/FloatingPragma/observer-patch-holography/releases/download/r1432/recovering_relativity_and_standard_model_structure_from_observer_overlap_consistency_compact.pdf.
- [33] B. Mueller, K. Xue, and K. A. Anirudha, “Reality as a Consensus Protocol,” GitHub release r1432 PDF, May 9, 2026. https://github.com/FloatingPragma/observer-patch-holography/releases/download/r1432/reality_as_consensus_protocol.pdf.
- [34] B. Mueller and K. Xue, “Federated Echoshedral Screen Microphysics,” OPH microphysics paper. https://github.com/FloatingPragma/observer-patch-holography/blob/main/paper/screen_microphysics_and_observer_synchronization.pdf.